

## Next milestone (29.07.)

Implement neural embeddings – either hardcode or softer version, using PyTorch

- Hardcode version:
  - No PyTorch, no ML libraries, **only numpy** (at least for the neural embeddings, can use PyTorch, etc. for GPT implementation)
    - Can use Counter and defaultdict from Collections
    - We can, but do not **have** to hardcode the optimiser (can use Adam, need to use at least SGD)
  - Measure perplexity
  - Implement early stopping (when validation error/loss diverges from training error to avoid overfitting to training set) with patience
    - Do not need to optimise for patience, but can
    - Save top k (the amount that fits reasonably on your disk) of model checkpoints (can name that file for validation score and iteration)
- Softer version:
  - Using PyTorch
  - Implement early stopping (when validation error/loss diverges from training error to avoid overfitting to training set) with patience
    - Do not have to optimise patience, but can
    - Save top k (the amount that fits reasonably on disk) of model checkpoints (can name that file for validation score and iteration)
  - Tune hyperparameters using a grid search **for each separately** and validation set (order is important: number of merges, learning rate, weights of interpolation) – do not have to do all of this to **pass**, but for 1.0
    - vocabulary size – gridsearch for max. 10 different amounts of merges
    - learning rate of optimiser
    - interpolation
  - Try versions with different optimisers
- Watch RAM during training, especially for higher batch sizes ( $\geq 32$ )